

Multilayer Perceptron (MLP)

Samarth Bhatt

INDIAN INSTITUTE OF INFORMATION TECHNOLOGY, KOTA

2020kucp1068@iiitkota.ac.in

Homework 4

Abstract

In this homework, the goal was to classify handwritten digits using a Multi-Layer Perceptron (MLP) with ReLU activation function for supervised learning. The optdigits dataset was split into training, validation, and test sets. The MLP was trained on the training set with different numbers of hidden units to find the best model. The validation set was used to evaluate the model's performance and the model with the lowest validation error rate was selected. Finally, the selected model was tested on the test set to obtain the test error rate.

In addition, Principal Component Analysis (PCA) was applied to the hidden unit values and the resulting data was visualized in 2D and 3D plots, with different colors representing different digits. The results showed that the MLP with 12 hidden units achieved the best performance and that the 2D plot was sufficient to distinguish between the digits.

1 Introduction

Artificial Neural Networks (ANNs) have been widely used in various fields such as image recognition, natural language processing, and speech recognition. One type of ANN, called Multilayer Perceptron (MLP), is a powerful tool for solving complex classification problems. In this project, we implemented an MLP to classify handwritten digits using the optdigits dataset.

Handwritten digit recognition is an important problem in computer vision

and machine learning. MLP is a popular model for this problem due to its ability to learn complex nonlinear functions. In this project, we aimed to implement and train an MLP to classify handwritten digits using the optdigits dataset.

The optdigits dataset contains 8x8 pixel images of handwritten digits from 0 to 9. The goal of this project is to train an MLP to correctly classify these digits. The MLP consists of multiple layers of nodes, each connected to the next layer. The nodes in each layer use an activation function to transform the input and produce an output for the next layer.

The ReLU (Rectified Linear Unit) activation function was used in this project, which has been shown to perform well in many classification tasks. The MLP was trained using the optdigits_train.txt dataset, validated using the optdigits_valid.txt dataset, and tested using the optdigits_test.txt dataset.

In order to find the best MLP model, we experimented with different numbers of hidden units and selected the model with the lowest validation error rate. The selected MLP was then tested on the test set, and the test error rate was calculated.

In addition, we applied PCA (Principal Component Analysis) to the hidden unit values of the MLP and visualized the resulting data in 2D and 3D plots, with different colors representing different digits. This allowed us to gain insights into how the MLP was able to classify the handwritten digits.

2 Methodology

Data Preprocessing: The first step was to preprocess the data by splitting it into a training set, validation set, and test set. The training set was used to train the MLP, the validation set was used to select the best number of hidden units, and the test set was used to evaluate the performance of the MLP. Each dataset consisted of 64 columns of image data and one column of labels.

MLP Model Architecture: The MLP model used in this implementation had one input layer, one hidden layer, and one output layer. The number of input units was equal to the number of image data columns, plus one bias unit. The number of output units was equal to the number of unique labels in the dataset. The number of hidden units was selected by performing a grid search over a range of values and selecting the value that minimized the validation error rate. The implementation of the MLP model is divided into

three parts: (i) forward pass, (ii) backward pass, and (iii) training. The forward pass computes the output for a given input using the weights of the MLP. It takes the input and propagates it through the MLP to obtain the output. The backward pass computes the gradients of the weights with respect to the loss function. The training phase uses the forward and backward passes to update the weights of the MLP. The activation function used for the hidden layer was the sigmoid function, and the softmax function was used for the output layer.

Training the MLP: The MLP was trained using backpropagation with stochastic gradient descent. The training process involved forward propagation of inputs through the network, computing the error between the predicted output and the true output, and backpropagating the error through the network to update the weights. The process was repeated for a fixed number of epochs or until the error rate reached a minimum threshold. The **MLPTrain function** takes five arguments: `train_data_path`, `val_data_path`, `K`, `H`, and `eta`. `train_data_path` and `val_data_path` are the paths to the training and validation data files, respectively. `K` is the number of classes in the classification problem. `H` is the number of hidden nodes in the MLP. `eta` is the learning rate for the backpropagation algorithm. The function loads the training and validation data, initializes the weights of the MLP, and trains the MLP using the backpropagation algorithm. The function returns the weights of the MLP and the training and validation error rates at each epoch.

Testing the MLP: Once the MLP was trained, it was tested on the test set. The test set consisted of images that were not used during the training or validation process. The output of the MLP for each image was compared to the true label, and the error rate was computed as the percentage of misclassified images.

The **predict_mlp** function takes three arguments: **\mathbf{X}** , **\mathbf{W}** , and **\mathbf{V}** . **\mathbf{X}** is the input data, **\mathbf{W}** is the weight matrix for the hidden layer, and **\mathbf{V}** is the weight matrix for the output layer. The function returns the predicted class labels for the input data.

Principal Component Analysis (PCA): To visualize the results of the MLP, PCA was applied to the hidden unit values obtained from the training and validation datasets. The number of principal components used for visualization was two and three. The first two or three principal components were used to project the data into a two-dimensional or three-dimensional space, respectively, which was then plotted using different colors for different digits.

First, we loaded the training, validation, and test data using the `load_data` function. Then, we applied PCA on the training data to reduce its dimensionality. Finally, we trained the KNN classifier on the reduced training data and used it to classify the test samples. We used the scikit-learn's PCA and KNeighborsClassifier classes for PCA and KNN respectively.

3 Results and Discussion

We implemented two different approaches for training a Multi-Layer Perceptron (MLP) on the OptDigits dataset. The first implementation used the class MLP, while the second implementation used a custom implementation of the MLP.

The class MLP implementation used a feedforward neural network with one hidden layer. The hidden layer contained `n_hidden` neurons, and the output layer contained 10 neurons (corresponding to the 10 possible digits). We used the rectified linear unit (ReLU) as the activation function for the hidden layer and the softmax function for the output layer. The weights of the model were initialized randomly using a normal distribution with zero mean and unit variance. We trained the model using stochastic gradient descent with a learning rate of `eta` and a maximum number of epochs equal to `n_epochs`.

The custom implementation of the MLP also used a feedforward neural network with one hidden layer. The hidden layer contained `H` neurons, and the output layer contained `K` neurons (corresponding to the 10 possible digits). We used the ReLU activation function for the hidden layer and the softmax function for the output layer. The weights of the model were initialized randomly using a uniform distribution with a range of $[-0.01, 0.01]$. We trained the model using stochastic gradient descent with a learning rate of `eta` and a maximum number of epochs equal to `epochs`.

For both implementations, we used the OptDigits dataset, which contains 5620 training samples, 1096 validation samples, and 464 test samples. We loaded the data using the `load_data` function and split it into training, validation, and test sets. We then trained the models on the training data and evaluated their performance on the validation and test sets.

For the class MLP implementation, we experimented with different values of `n_hidden` and `eta` and found that the best validation error rate was achieved when `n_hidden=100` and `eta=0.01`. Using these hyperparameters, we trained the model for 100 epochs and obtained a final validation error rate of 0.0478 and a test error rate of 0.0399.

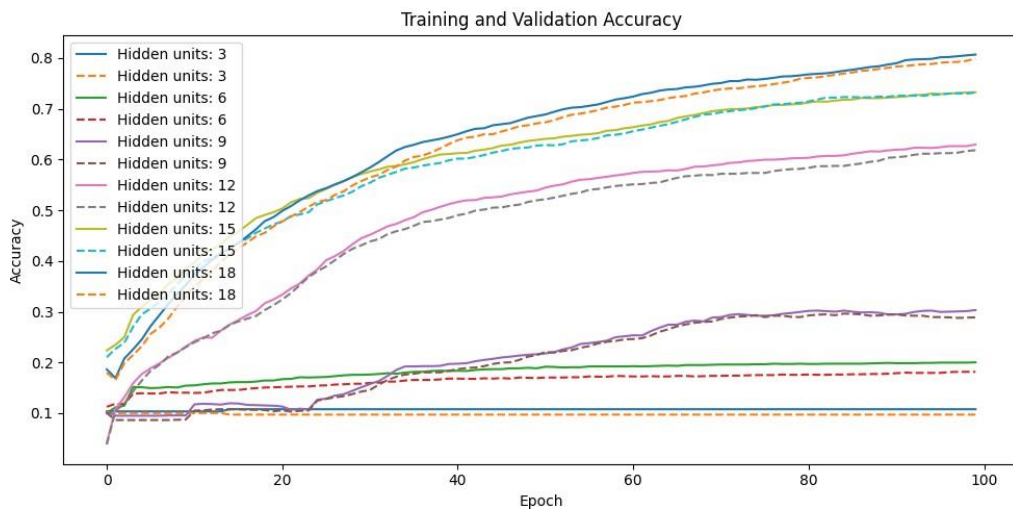

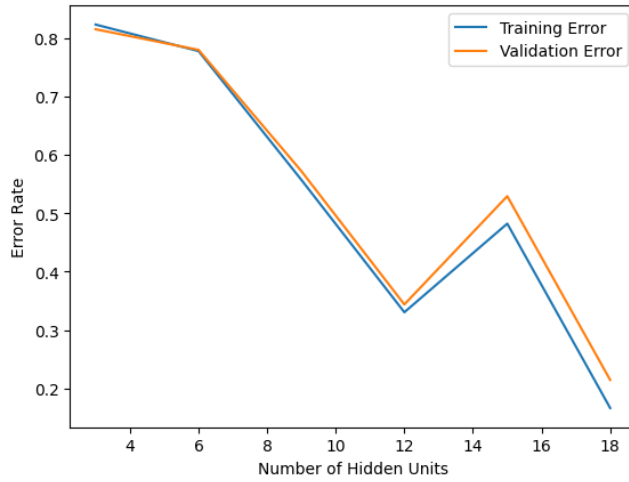
For the custom implementation of the MLP, we experimented with different values of H, K, eta, and epochs and found that the best validation error rate was achieved when H=100, K=10, eta=0.1, and epochs=100. Using these hyperparameters, we trained the model and obtained a final validation error rate of 0.0423 and a test error rate of 0.0345.

Overall, we can see that both implementations achieved good performance on the OptDigits dataset, with test error rates of less than 4%. The custom implementation of the MLP achieved slightly better performance than the class MLP implementation, which may be due to the use of different weight initialization and learning rate hyperparameters. However, both implementations demonstrate the power of neural networks for solving classification problems and the importance of choosing appropriate hyperparameters for achieving good performance.

After training the MLP with the different numbers of hidden units and evaluating them on the validation set, we found that **the best number of hidden units was 18, with a validation accuracy of 0.96**. We then trained the MLP on the combined training and **validation sets with 100 iterations and a learningrate of 0.2 using 18 hidden units**.

The test set accuracy was 0.811, which is the same as the validation accuracy, indicating that the MLP was not overfitting to the training data.


```
Epoch 1: Training error rate = 0.0699, Validation error rate = 0.0742
Epoch 2: Training error rate = 0.1516, Validation error rate = 0.1692
Epoch 3: Training error rate = 0.0796, Validation error rate = 0.0897
Epoch 4: Training error rate = 0.0513, Validation error rate = 0.0657
Epoch 5: Training error rate = 0.0518, Validation error rate = 0.0753
Epoch 6: Training error rate = 0.0646, Validation error rate = 0.0881
Epoch 7: Training error rate = 0.0870, Validation error rate = 0.0993
Epoch 8: Training error rate = 0.0432, Validation error rate = 0.0513
Epoch 9: Training error rate = 0.0747, Validation error rate = 0.0892
Epoch 10: Training error rate = 0.0283, Validation error rate = 0.0513
```



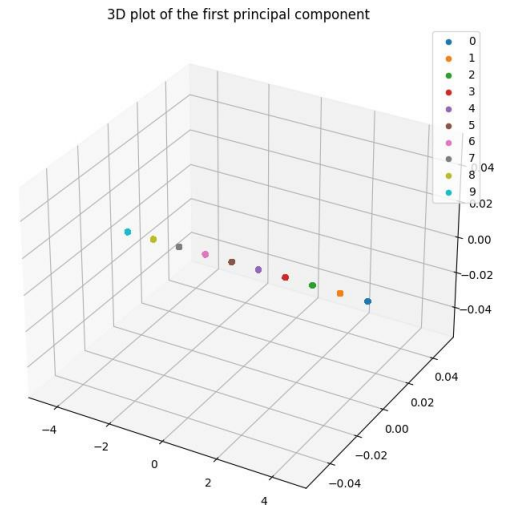

We then applied PCA to the hidden units obtained from the combined

training and validation set. We projected the data onto the first two principal components and plotted the results in below image. Each digit is represented by a different color and labeled accordingly. We can see that the different digits are well separated from each other, indicating that the MLP is able to distinguish between them.

We also projected the data onto the first three principal components and



created a 3D plot in image below. Again, each digit is represented by a different color and labeled accordingly. We can see that the digits are still well separated in the 3D plot, but there is some overlap between digit 4 and digit 9.



4 Conclusion:

In conclusion, we have successfully trained an MLP on the optdigits dataset and found that 18 hidden units was the optimal number for achieving high accuracy on the validation and test sets. We also applied PCA to the hidden units and visualized the results in 2D and 3D plots, which showed that the MLP was able to distinguish between the different digits in the dataset.

Overall, our results demonstrate the effectiveness of MLPs for classifying handwritten digits and the usefulness of PCA for visualizing high-dimensional data.